

Supplementary Material for: Robust Branch-Safe Participation for Sequential Mobile Crowdsensing

Lili He^a, Sheng Jiang^a, Linghui Lyu^a, Lei Zhang^{a,*}

^a*College of Information and Electronic Technology, Jiamusi University, Jiamusi 154007, China*

Abstract

This supplement records implementation details that are useful for verification but not necessary for following the main argument: comparator adaptations, finite-model robustness and boundary tests, construction of the data-driven ambiguity set, and reproducibility information. It does not introduce additional primary claims. The post-hoc status of the replay extensions and the finite-model scope of the robustness guarantee are preserved.

1. Comparator adaptations and validation

For the pointwise-maximal-leakage comparator, let $l_i = \Pr(Y = 1 \mid X = i)$ and $m = \sum_i b_i l_i$, with $0 \leq l_i \leq \alpha \mathbf{1}\{i \in S\}$. The task-channel specialization maximizes report probability subject to the published pointwise leakage constraints $l_i \leq e^\epsilon m$ and $1 - l_i \leq e^\epsilon (1 - m)$ for every state with positive prior mass [1]. Randomized instances are checked against a general linear program. This is a task-channel specialization rather than a reproduction of every component of the original application.

PRIVIC learns a noisy location channel through Blahut–Arimoto and incremental Bayesian updates [2]. Our BA/IBU/GIBU updates are checked against the authors’ public implementation. The task adapter emits a truthful report only when

*Corresponding author.

Email address: 8213662@163.com (Lei Zhang)

the true state is eligible and the sampled location-channel outcome satisfies the task predicate. Because the adapter changes the output semantics, PRIVIC-T does not inherit the original geo-indistinguishability guarantee. Offline fitting is excluded from online gate timing.

The Bayesian-privacy comparator constrains each positive-probability output by $D_{\text{KL}}(b^v||b) \leq \kappa$ [3]. The successful-report divergence is determined by task support, while the silence bound is solved by bisection. It is used as a per-output diagnostic rather than as a cumulative session guarantee.

2. Finite-model robustness and boundary characterization

The hand-grid ambiguity set uses $v \in \{0.05, 0.3, 0.8\}$ and $\alpha \in \{0.3, 0.648, 0.9\}$. Twenty new synthetic seeds, four users per seed, and 48 slots yield 80 trajectories per reported row at a fixed operating point. When the informed attacker’s exact generating model is contained in the set, R-BSP records no local-cap violation in the reported in-set conditions, as expected from the finite-set construction, while retaining about 39.7% of externally completable normal-task opportunities.

Table 1: Finite-model R-BSP stress at $\rho = 0.1$.

Client model	BSP viol.	R-BSP viol.	BSP retention	R-BSP retention
Nominal	0.0%	0.0%	50.9%	39.7%
$\widehat{\alpha} = 0.3$	12.6%	0.0%	54.2%	39.7%
$\widehat{\alpha} = 0.9$	2.6%	0.0%	47.3%	39.7%
$\widehat{v} = 0.05$	2.8%	0.0%	46.9%	39.7%
$\widehat{v} = 0.8$	14.6%	0.0%	55.2%	39.7%

The same ambiguity set is then evaluated against correct attacker models that are not exact members. The outside-set value $\alpha = 0.98$ yields a 6.93% local-cap

violation rate with a cluster-bootstrap 95% interval of 5.70–8.31%. Interior unlisted values yield zero or rare violations in this finite sample. These results support only a finite-model guarantee and are not evidence of robustness to arbitrary distribution shift.

Table 2: Boundary characterization for attacker models not explicitly enumerated in the hand-grid ambiguity set.

Attacker model	BSP viol.	R-BSP viol.	BSP ret.	R-BSP ret.
$\alpha = 0.50$ (interior)	1.15%	0.00%	52.6%	36.3%
$\alpha = 0.98$ (outside)	18.33%	6.93%	50.0%	40.4%
$v = 0.50$ (interior)	0.65%	0.05%	54.4%	36.4%
$v = 0.95$ (outside)	0.65%	0.03%	55.6%	35.4%

3. Data-driven ambiguity construction

The replay ambiguity set is constructed without using the 53 isolated test windows. The 23 development windows define a coarsened reflecting-walk likelihood

$$\Pr(\text{change} \mid i) = vd_i/4,$$

where d_i is the number of in-grid neighbors. Only change versus no-change events enter this likelihood; destination direction and jump distance are not fitted. Two thousand development state-path-group bootstrap replicates generate mobility candidates, and 2,000 validation state-path-group bootstrap replicates select the smallest set covering at least 95% of predictive-likelihood winners.

The validation winner distribution is concentrated on two development-only smoothed population-prior candidates, with selection frequencies 84.15% and 12.80%. Retaining these two models gives 96.95% cumulative validation selection mass, above the prespecified 95% target. The development MLE is 0.0514 with

a 95% bootstrap interval 0.0298–0.0754; the validation point estimate is 0.0817. Availability remains externally specified because GeoLife contains trajectories rather than task-delivery, willingness, or deadline observations.

A validation-best single-model BSP control is retained to separate improved model fit from finite-set robustness. The strict data-driven R-BSP and the disclosure-floor-aware RDF-BSP service frontier are reported in the main manuscript rather than duplicated here.

4. Reproducibility record

The repository stores deterministic preprocessing, split metadata, configuration files, seeds, trajectory-level summaries, bootstrap outputs, provenance records, analysis code, and manuscript sources. The effective GeoLife analysis split contains 23 development, 12 validation, and 53 test windows after model-input-group isolation. User IDs and complete 48-slot discrete state-path fingerprints are disjoint across the effective splits.

GeoLife uncertainty intervals resample complete state-path groups rather than treating slots or duplicate model-input paths as independent replicates. The current automated suite contains 30 tests covering nominal, robust, and disclosure-floor posterior branches; gate reductions and maximality; informed-attacker integration; transition conservation; GeoLife split and missing-minute handling; numerical replay; PML linear-program checks; PRIVIC author-code checks; and policy-envelope optimization.

The data-driven replay extension remains labelled post-hoc because the test replay had been inspected in earlier project stages. No mobile-device energy measurement, production MCS task-log study, or independent external reproduction is claimed. The original Microsoft GeoLife archive and third-party article full texts

are not redistributed in the submission package.

Funding

This work was supported in part by the Heilongjiang Provincial University Teacher Scientific Research Business Fee under Grant 2018-KYYWF0941, in part by the Heilongjiang Provincial Natural Science Foundation Joint Fund Cultivation Project under Grant PL2024F002, in part by the National Foreign Expert Key Support Project under Grant D2O25O185, and in part by the Heilongjiang Provincial Undergraduate University Excellent Young Teacher Fundamental Research Support Plan under Grant YQJH2024239.

References

- [1] L. Grosse, S. Saeidian, T. J. Oechtering, Extremal Mechanisms for Pointwise Maximal Leakage, *IEEE Transactions on Information Forensics and Security* 19 (2024) 7952–7967. doi:10.1109/tifs.2024.3449556.
- [2] S. Biswas, C. Palamidessi, PRIVIC: A privacy-preserving method for incremental collection of location data, *Proceedings on Privacy Enhancing Technologies* 2024 (1) (2024) 582–596. doi:10.56553/popets-2024-0033.
- [3] R. Eilat, K. Eliaz, X. Mu, Bayesian privacy, *Theoretical Economics* 16 (4) (2021) 1557–1603. doi:10.3982/te4390.